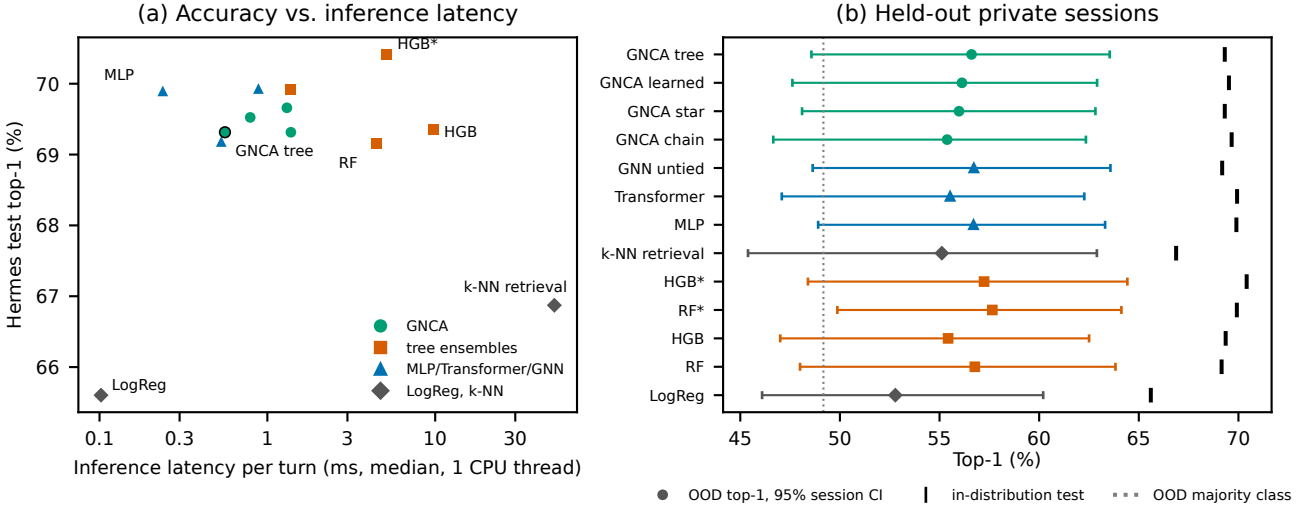
(a) Accuracy vs. inference latency
Hermes test top-1 (%)
70
69
68
67
66
0.1
0.3
1
3
10
30
Inference latency per turn (ms, median, 1 CPU thread)
HGB*
MLP
HGB
GNCA tree
RF
k-NN retrieval
LogReg
GNCA
tree ensembles
MLP/Transformer/GNN
LogReg, k-NN
(b) Held-out private sessions
GNCA tree
GNCA learned
GNCA star
GNCA chain
GNN untied
Transformer
MLP
k-NN retrieval
HGB*
RF*
HGB
RF
LogReg
45
50
55
60
65
70
Top-1 (%)
OOD top-1, 95% session CI
in-distribution test
OOD majority class